

Optical character recognition

Виконав


Левадний Микола, МІ-4



Dataset

A-Z Handwritten Alphabets in .csv format

The dataset contains 26 folders (A-Z) containing handwritten images in size 28x28 pixels, each alphabet in the image is centre fitted to 20x20 pixel box.




Add data augmentation

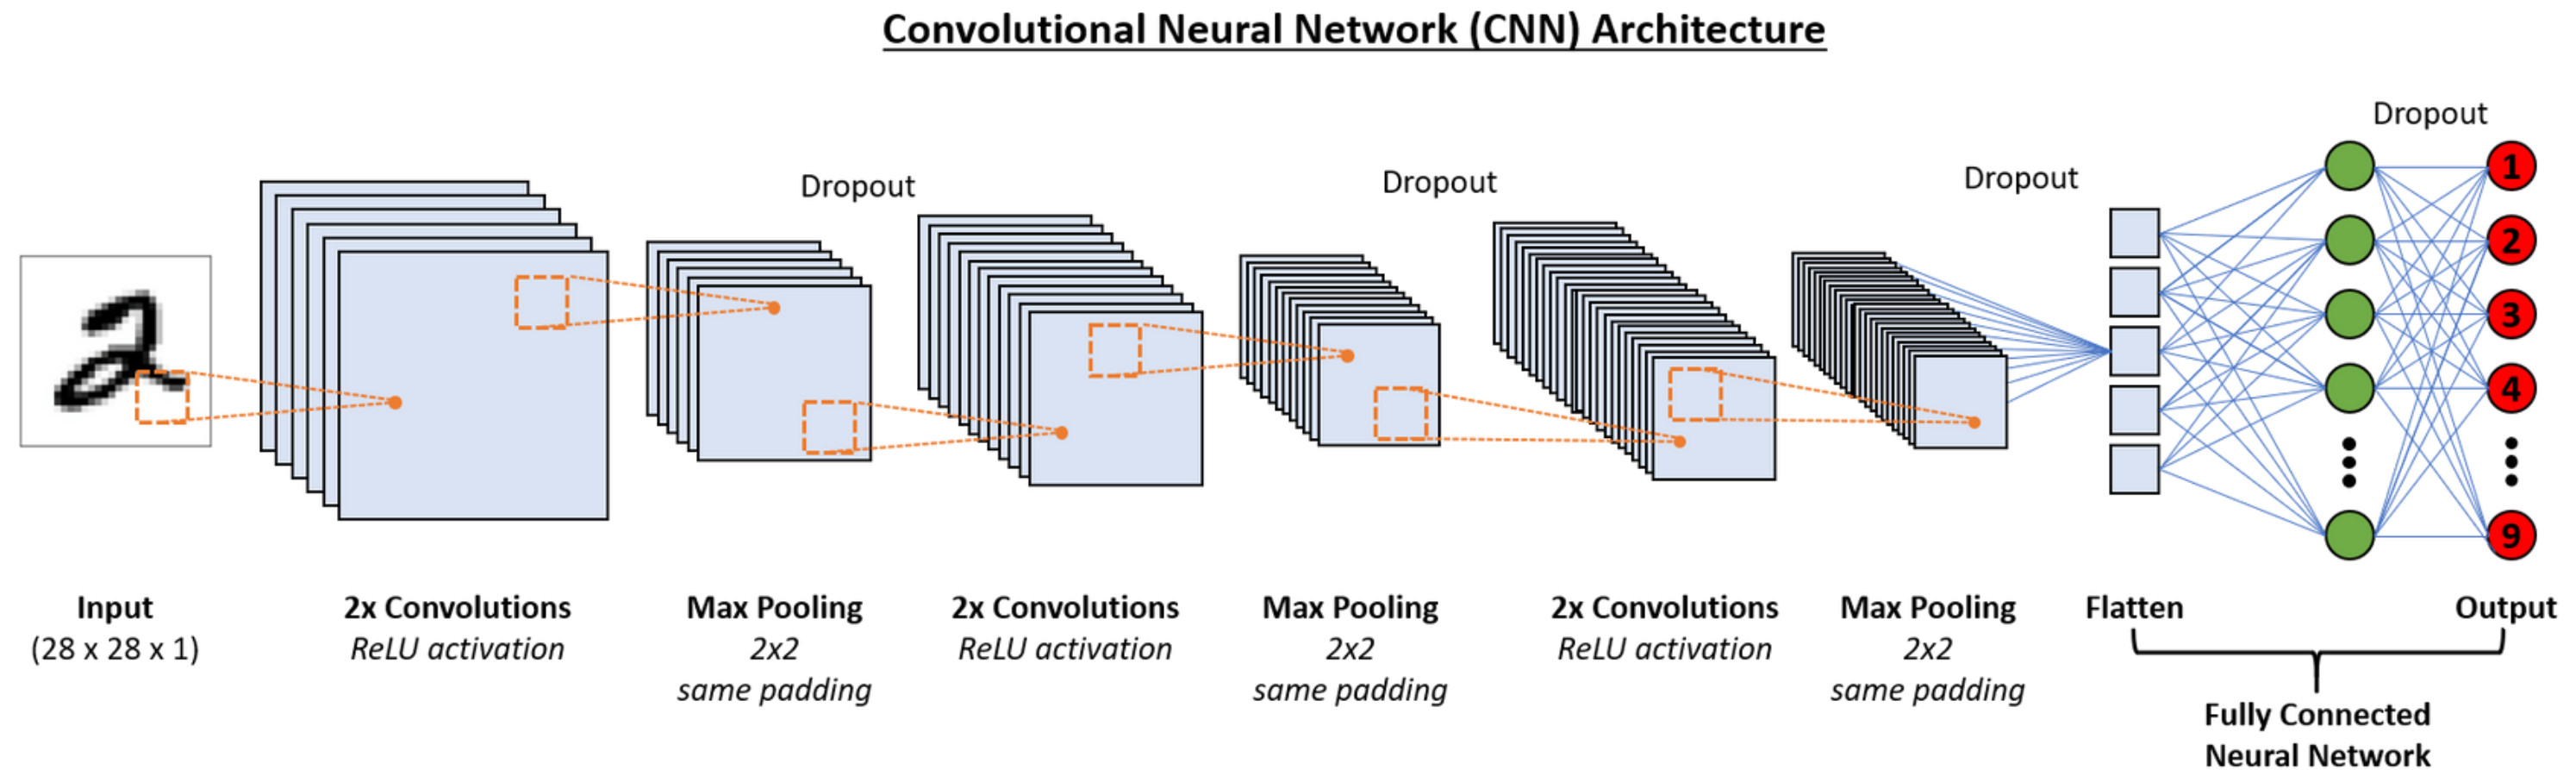
Augmentation allows you to expand the dataset with the help of already received data

- Small rotation (prevent letter overturn)
- Changing contrast of image
- Zooming

```
data_augmentation = tf.keras.Sequential([  
    layers.experimental.preprocessing.RandomRotation(0.08),  
    layers.experimental.preprocessing.RandomContrast(0.12),  
    layers.experimental.preprocessing.RandomZoom(height_factor=(0.05,0.1))  
])
```



CNN model



A Convolutional Neural Network, also known as CNN or ConvNet, is a class of neural networks that specializes in processing data that has a grid-like topology, such as an image. A digital image is a binary representation of visual data. It contains a series of pixels arranged in a grid-like fashion that contains pixel values to denote how bright and what color each pixel should be.

CNN realization code snippet

(using keras)

```
[39] model = Sequential()
model.add(data_augmentation)
#CNN
# input -> conv -> maxpool -> conv -> maxpool .....->flattened vector->
# .                hidden layer -> hidden layer -> softmax layer
model.add(Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(28,28,1)))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))

model.add(Conv2D(filters=64, kernel_size=(3, 3), activation='relu', padding = 'same'))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))

model.add(Conv2D(filters=128, kernel_size=(3, 3), activation='relu', padding = 'valid'))
model.add(MaxPool2D(pool_size=(2, 2), strides=2))

model.add(Flatten())

model.add(Dense(64,activation ="relu"))
model.add(Dense(128,activation ="relu"))

model.add(Dense(26,activation ="softmax"))
```

model.summary()

Model: "sequential_10"

| Layer (type) | Output Shape | Param # |
|---------------------------------|--------------------|---------|
| ===== | | |
| sequential_9 (Sequential) | (None, 28, 28, 1) | 0 |
| conv2d_20 (Conv2D) | (None, 26, 26, 32) | 320 |
| max_pooling2d_19 (MaxPooling2D) | (None, 13, 13, 32) | 0 |
| conv2d_21 (Conv2D) | (None, 13, 13, 64) | 18496 |
| max_pooling2d_20 (MaxPooling2D) | (None, 6, 6, 64) | 0 |
| conv2d_22 (Conv2D) | (None, 4, 4, 128) | 73856 |
| max_pooling2d_21 (MaxPooling2D) | (None, 2, 2, 128) | 0 |
| flatten_7 (Flatten) | (None, 512) | 0 |
| dense_20 (Dense) | (None, 64) | 32832 |
| dense_21 (Dense) | (None, 128) | 8320 |
| dense_22 (Dense) | (None, 26) | 3354 |

=====

Total params: 137,178
Trainable params: 137,178
Non-trainable params: 0

=====

Training results

Compiling the Model

```
✓ [61] model.compile(optimizer = Adam(learning_rate=0.001), loss='categorical_crossentropy', metrics=['accuracy'])
```

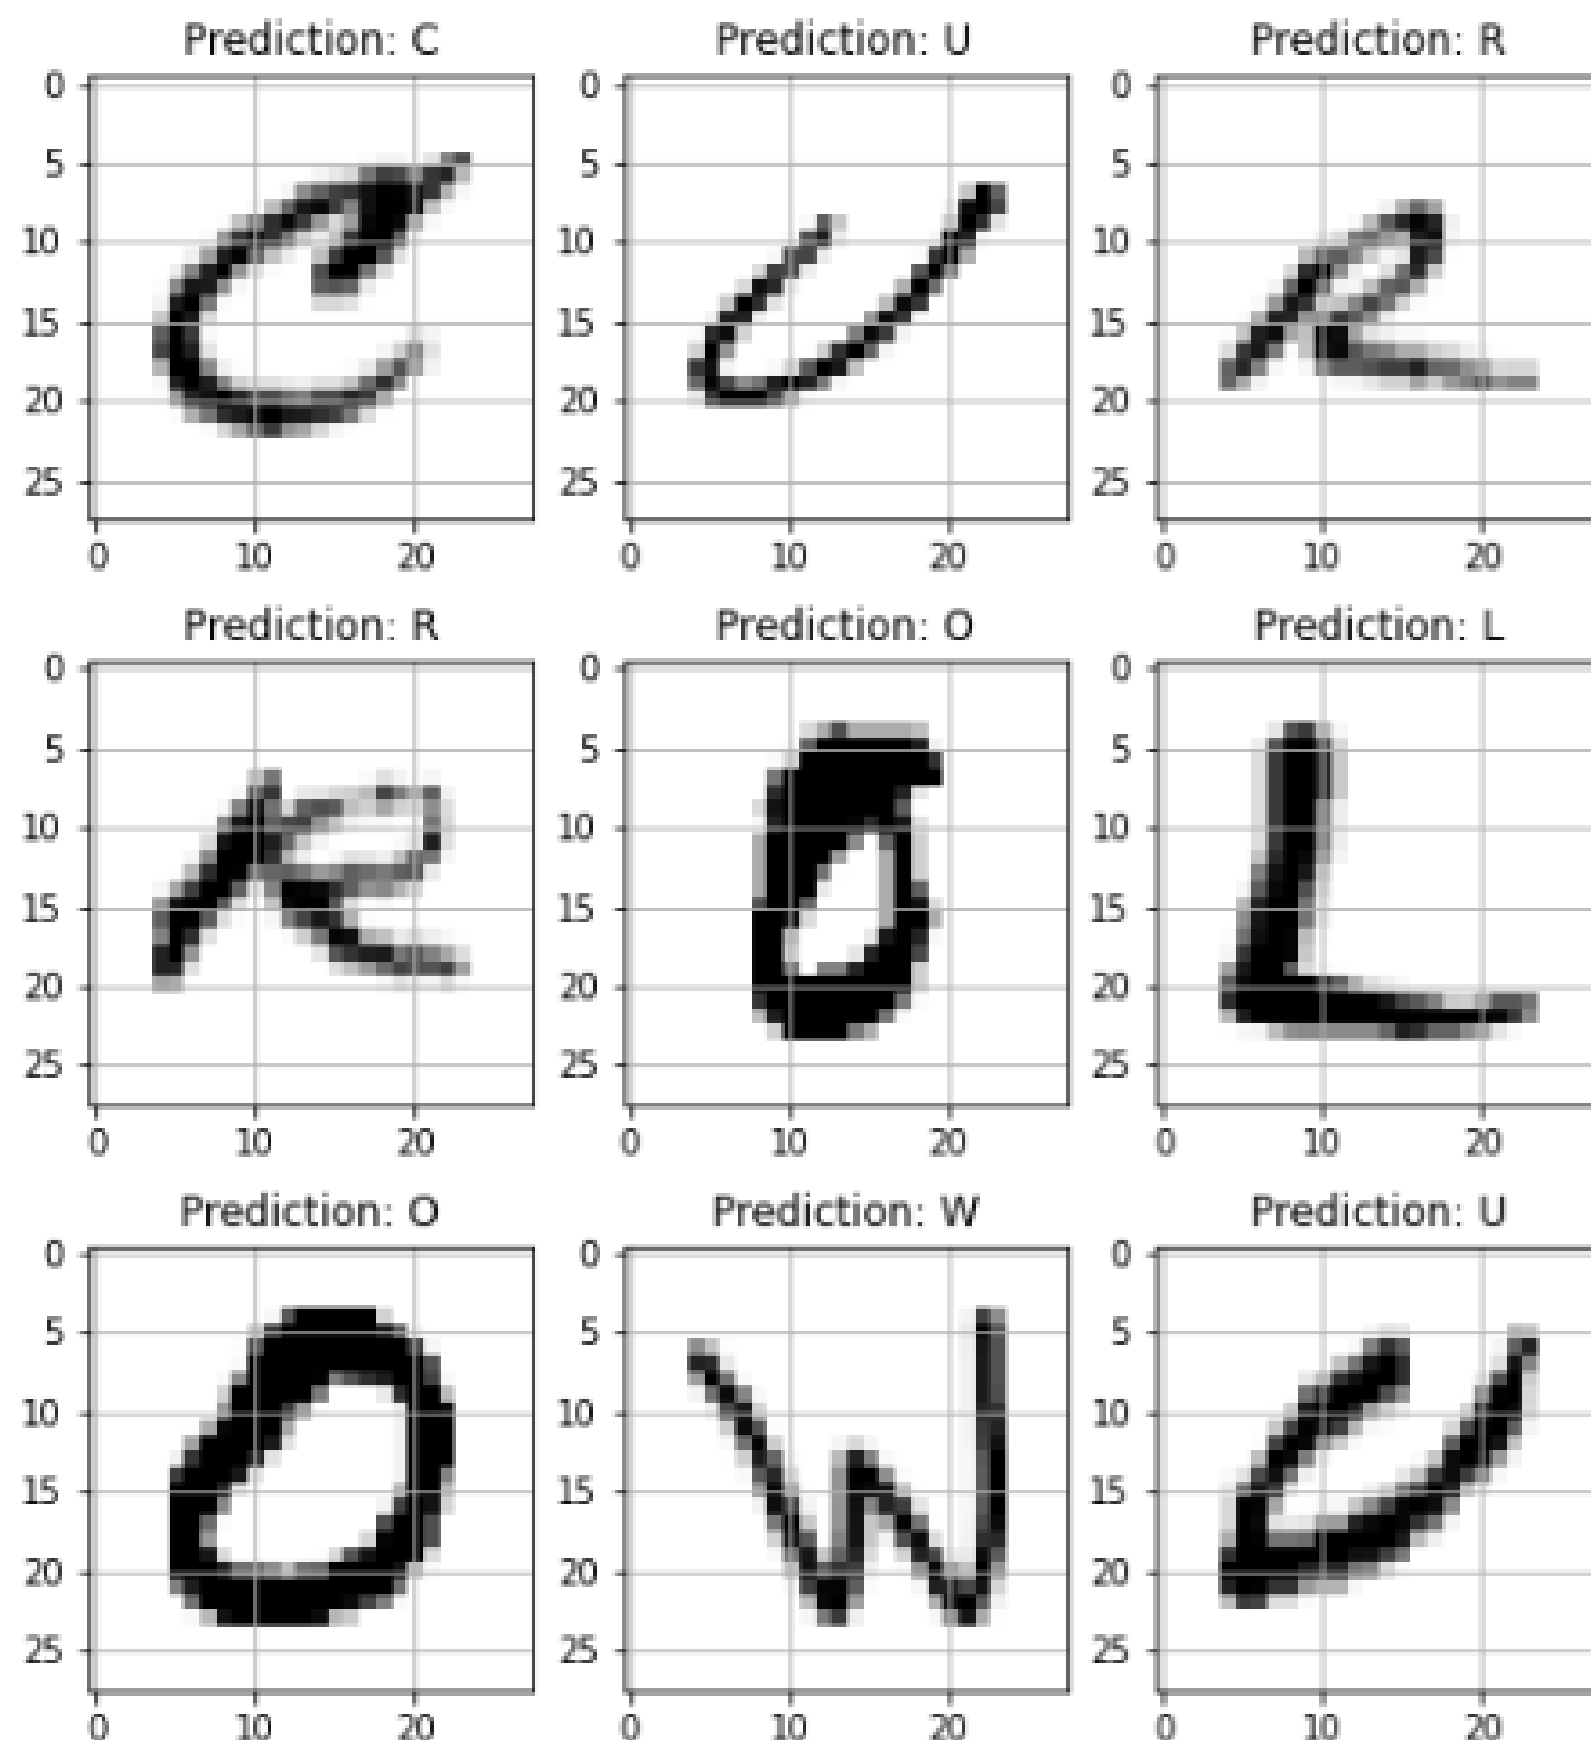
Starting the Training

```
✓ [62] history = model.fit(X_train, train_yOHE, epochs=10, validation_data = (X_test, test_yOHE))
```

```
Epoch 1/10
9312/9312 [=====] - 40s 4ms/step - loss: 0.1580 - accuracy: 0.9564 - val_loss: 0.0951 - val_accuracy: 0.9725
Epoch 2/10
9312/9312 [=====] - 38s 4ms/step - loss: 0.0713 - accuracy: 0.9801 - val_loss: 0.0677 - val_accuracy: 0.9822
Epoch 3/10
9312/9312 [=====] - 37s 4ms/step - loss: 0.0597 - accuracy: 0.9835 - val_loss: 0.0753 - val_accuracy: 0.9810
Epoch 4/10
9312/9312 [=====] - 37s 4ms/step - loss: 0.0557 - accuracy: 0.9850 - val_loss: 0.0650 - val_accuracy: 0.9839
Epoch 5/10
9312/9312 [=====] - 43s 5ms/step - loss: 0.0513 - accuracy: 0.9859 - val_loss: 0.0662 - val_accuracy: 0.9824
Epoch 6/10
9312/9312 [=====] - 38s 4ms/step - loss: 0.0503 - accuracy: 0.9867 - val_loss: 0.0591 - val_accuracy: 0.9853
Epoch 7/10
9312/9312 [=====] - 38s 4ms/step - loss: 0.0491 - accuracy: 0.9872 - val_loss: 0.0655 - val_accuracy: 0.9846
Epoch 8/10
9312/9312 [=====] - 37s 4ms/step - loss: 0.0498 - accuracy: 0.9874 - val_loss: 0.0738 - val_accuracy: 0.9833
Epoch 9/10
9312/9312 [=====] - 37s 4ms/step - loss: 0.0468 - accuracy: 0.9884 - val_loss: 0.0680 - val_accuracy: 0.9856
Epoch 10/10
9312/9312 [=====] - 37s 4ms/step - loss: 0.0501 - accuracy: 0.9880 - val_loss: 0.0642 - val_accuracy: 0.9857
```

Test on validation set

Accuracy result: 0.9857





Making Predictions


```
[63] model.evaluate(X_test, np.array(test_yOHE))
```

```
2328/2328 [=====] - 6s 3ms/step - loss: 0.0642 - accuracy: 0.9857  
[0.06420440971851349, 0.9857430458068848]
```

Training and validation accuracy and loss trend




Confusion Matrix



Typical errors

We can notice that the model is trained well enough, that's why there are not many errors. But still there are enough cases where there were mix-ups O-D and N-H





**Thanks for
attention!**